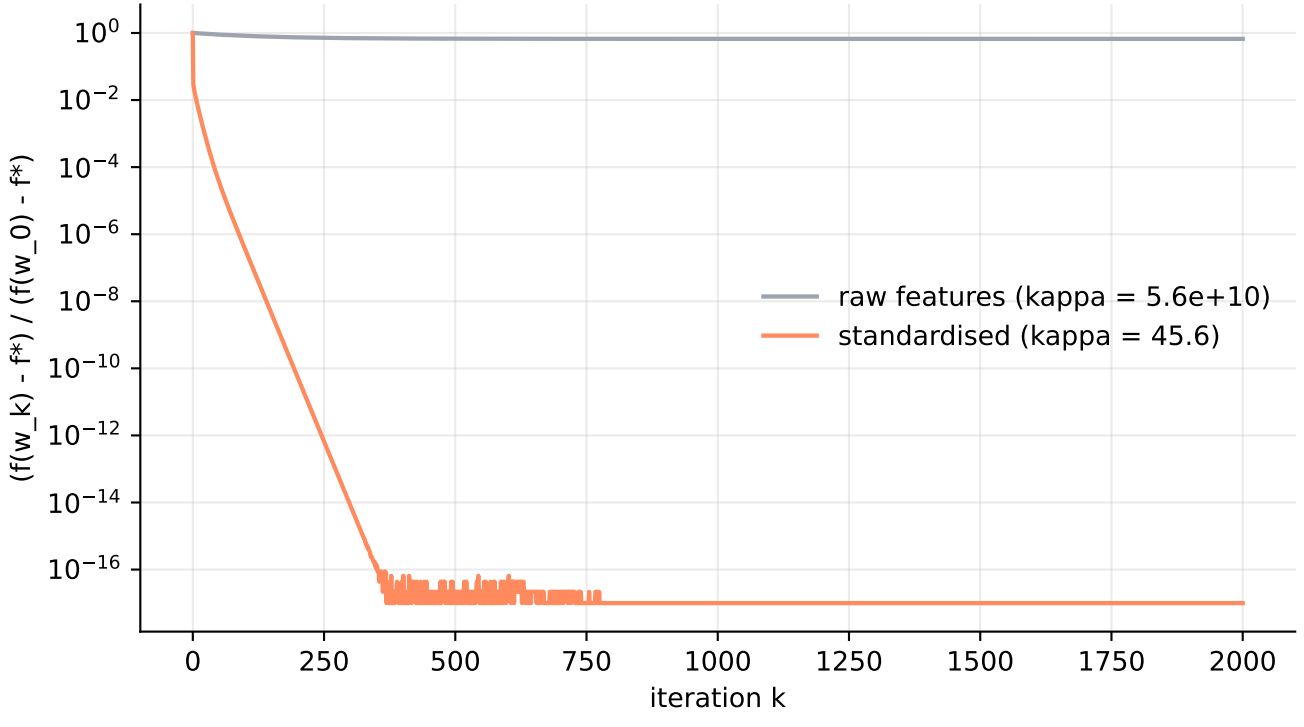
Same algorithm, same optimal step: scaling changes everything
(f(w_k) - f*) / (f(w_0) - f*)
10^{0}
10^{-2}
10^{-4}
10^{-6}
10^{-8}
10^{-10}
10^{-12}
10^{-14}
10^{-16}
raw features (kappa = 5.6e+10)
standardised (kappa = 45.6)
0
250
500
750
1000
1250
1500
1750
2000
iteration k